

Building a Business Case for an MCP-Enabled Eligibility-Screening Tool for Make-A-Wish Healthcare Partners

TL;DR

- **The opportunity is real and timed correctly.** Make-A-Wish America granted a record 16,295 wishes in FY24 but still reaches only roughly one in two eligible children — and U.S. hospitals are crossing the 50% threshold for generative-AI adoption in 2025, with pediatric AMCs (Boston Children's, CHOP, Stanford Children's, Cedars-Sinai) already named as launch partners for ChatGPT for Healthcare. An MCP integration that surfaces likely-eligible wish kids inside a clinician's existing AI workspace lands in a market that is buying, not browsing. (wish)
 - **The technical and legal scaffolding now exists.** Anthropic's HIPAA-ready Claude Enterprise (Dec 2025), OpenAI's ChatGPT for Healthcare with BAA (Jan 2026), Epic's Art/Penny/Emmie agents, and the Model Context Protocol (Anthropic, Nov 2024) collectively make a HIPAA-eligible, EHR-aware screening agent feasible without bespoke per-EHR builds — provided BAAs, ONC HTI-1 transparency, and PHI redaction at the prompt layer are in place.
 - **Pitch it to children's-hospital social workers, oncology care coordinators, and child-life teams — not IT.** Make-A-Wish's approved referrers are providers, nurses, social workers, and child-life specialists; the bottleneck is recognition and time, not data. A 30-second MCP agent that scans the inpatient roster for the 12 qualifying specialties (oncology, cardiology, neurology, hematology, genetics, pulmonology, immunology, endocrinology, rheumatology, nephrology, ID, GI), checks age 2½-17, and drafts a pre-filled referral inquiry to wish.org is the wedge. (Wish)
-

Key Findings

1. Make-A-Wish eligibility and referral mechanics

- **Eligibility criteria (national, per wish.org/refer-a-child and worldwish.org):**
 - Age 2½ to <18 at time of referral (worldwide criterion most commonly stated as 3-17, but the U.S. site uses 2½ as the floor; wishes for children under 3 sometimes require board approval).
 - **Critical illness** defined as "a progressive, degenerative, or malignant condition that has placed the child's life in jeopardy." Not required to be terminal. (Wish)
- (Wishmakersball)

- **No prior wish** granted by Make-A-Wish or any other wish-granting organization. (Wish)
- **Financial status is never considered.** (Wishmakersballjax) (Wishmakersball)
- **Approved referral sources (only these four can initiate an official referral):**
 1. Medical professionals treating the child — doctors, nurse practitioners, PAs, RNs, social workers, child-life specialists (Wish) (Wish)
 2. The child's parent or legal guardian
 3. Family members with detailed knowledge of the child's current medical condition (Wish)
 4. The potential wish child themselves
- **12 qualifying specialties** (Make-A-Wish International's published list, used by U.S. chapters): Neurology (SMA, Rett, intractable epilepsy), Cardiology (complex congenital defects, transplant, heart failure), Hematology (sickle cell, aplastic anemia, SCID), Genetics (trisomy 13/18, skeletal dysplasias, metabolic syndromes), Oncology (leukemia, brain tumors, stem-cell transplant), Pulmonology (CF with complications, ventilator dependence), Immunology, Endocrinology, Rheumatology, Nephrology, Infectious Disease, Gastroenterology. **Cystic fibrosis no longer automatically qualifies as of Jan 1, 2024** — it now requires additional complicating factors (per the National Medical Advisory Council update). (Make-A-Wish + 3)
- **Medical professional of record:** The child's **treating physician confirms medical eligibility** and signs off on safety/timing of the wish. The **Make-A-Wish National Medical Advisory Council (19 volunteer physicians)** plus **200+ chapter-specific medical advisors** review borderline cases. (Make-A-Wish + 2)
- **The under-referral problem (the business case in one number):** Make-A-Wish has long stated it reaches **only one in two eligible U.S. children** (wish.org/wishes-forever). Its FY24 framing is "for every wish we fulfill, two more are waiting." The cited cause, in Make-A-Wish's own words via Contemporary Pediatrics, is **provider misunderstanding** — clinicians wrongly believe the program is only for terminal kids, so children in oncology remission, transplant recovery, or stable-but-degenerative neurology don't get referred. (Wish + 2)
- **Cited annual incidence:** Make-A-Wish itself states ~27,000 U.S. children are diagnosed with a critical illness each year ("three children every hour"). FY24 wishes granted enterprise-wide: **16,295** (FY24 Annual Report, p.3 — "a historic milestone ... more than ever throughout our organization's 45-year history"). FY24 total revenue \$534.3M, expenses \$500.6M, program services 79.2% of operating spend. (Wish)

2. AI adoption in U.S. healthcare — current state (2024-2026)

- **Adoption is past the tipping point.** A JAMA Network Open survey (Everson, Nong & Richwine; ONC + Univ. of Minnesota; n=2,174 nonfederal acute-care hospitals) published

Dec 1, 2025 found:

- **31.5% of U.S. hospitals already using generative AI integrated with the EHR in 2024** (PubMed)
- **24.7% planning adoption within a year ("fast followers")** (PubMed)
- **43.7% delayed adopters** (no plans / 5+ year horizon / unsure) (PubMed)
- System-affiliated, teaching hospitals, and Epic-EHR hospitals adopt fastest; independent hospitals and high-Medicaid-discharge hospitals lag. (Beckers Hospital Review)
- **McKinsey Q4 2025 healthcare survey:** 50% of U.S. healthcare organizations have implemented generative AI by end of 2025 (up from 25% in late 2023 and 47% in 2024). >80% of leaders report at least one production GenAI use case; 100% report plans to pursue more. (Beckers Hospital Review)
- **Menlo Ventures 2025 State of AI in Healthcare (released Oct 21, 2025; n=700+ senior healthcare executives): 85% of GenAI spend in healthcare flows to startups, not incumbents.** In the \$600M ambient-scribe category, the market splits Nuance DAX Copilot 33%, Abridge 30%, Ambience 13%, Suki 10%, and others 14% — meaning **Abridge and Ambience together captured "nearly 70% of the new [GenAI-native] market"** even though Microsoft Nuance DAX had previously been deployed in 77% of U.S. hospitals. (Menlo Ventures + 2)
- **Scottsdale Institute 2024 survey (43 leading U.S. health systems, 64% response):** 100% of respondents reported adoption activity in ambient documentation; >50% reported "a high degree of success" with clinical documentation AI specifically. Imaging AI: only 19% report high success. (PubMed Central)
- **AMA 2024:** 2 in 3 physicians using health AI; 21–31% of physicians/family physicians use AI for clerical support (message drafting, ambient documentation).

3. Children's hospitals leading AI

Direct, named pediatric-system AI leaders documented in the trade press and primary sources:

- **Boston Children's Hospital** — Chief Innovation Officer **John Brownstein** runs the AI strategy; hired what he calls "unicorns" including **Dinesh Rai, MD** as inaugural prompt engineer with medical expertise. Piloting Red Hat OpenShift AI in pediatric radiology (announced May 21, 2024). Named launch partner for **ChatGPT for Healthcare** (OpenAI, Jan 2026) alongside AdventHealth, Baylor Scott & White, Cedars-Sinai, HCA, Memorial Sloan Kettering, Stanford Medicine Children's Health, and UCSF. (Northeastern + 2)
- **Stanford Medicine Children's Health** — ChatGPT for Healthcare launch partner; Stanford Health Care has built its own EHR-integrated chatbot called **ChatEHR** under Dr. Nigam Shah, chief data scientist. (Beckers Hospital Review)
- **Children's Hospital of Philadelphia (CHOP), Texas Children's, Children's Hospital of Colorado, Lurie Children's, Seattle Children's, Akron Children's, Dayton Children's,**

UPMC Children's — all named as rapid adopters of the **Abridge ambient-AI scribe platform** (Abridge case study + Becker's interviews). (Abridge)

- **Texas Children's** previously deployed DRUID conversational AI for COVID-era vaccination verification — demonstrating willingness to deploy chat AI at scale.
- **St. Jude** is a leading research center but is less publicly profiled as a clinical generative-AI deployer relative to peers.

4. The dominant AI tools entering hospital workflows

- **OpenAI ChatGPT for Healthcare (launched Jan 2026)**: Built on GPT-5 with physician-led evaluation. BAA-eligible. Launch customers include Boston Children's and Stanford Children's. Templates for discharge summaries, prior-auth, clinical letters, differential diagnoses. Integrates with Microsoft SharePoint. PHI remains under organization control with data residency, audit logs, customer-managed encryption keys. (Fierce Healthcare + 3)
- **Anthropic Claude — HIPAA-ready Enterprise (Dec 2025)**: Click-to-accept BAA enabled from organization settings. Used by **Banner Health** (33-hospital Arizona system) which built "BannerWise" on Claude — reducing oncology chart summarization from ~8 hours per patient to automated, with 80–85% of users reporting time savings. AWS Bedrock and Google Vertex AI also offer HIPAA-eligible Claude access. **Claude Code, Cowork, and beta features are NOT BAA-covered.** (Claude + 3)
- **Microsoft 365 Copilot + Azure OpenAI**: HIPAA-supportable inside Microsoft Cloud for Healthcare with BAA. **Microsoft Dragon Copilot** (formerly Nuance DAX) integrated into Epic AI Charting. Microsoft launched a consumer "Copilot Health" tab at HIMSS 2026, but the enterprise stack is the relevant lane for hospital deployments. (Healthcare Dive)
- **Google Gemini in Workspace** is HIPAA-supportable with a BAA — but Gemini-in-Chrome and NotebookLM are **excluded from the BAA.** (John D. Cook)
- **Epic AI agents (Art / Emmie / Penny)**: As of early 2026, **85% of Epic's healthcare customers are live with at least one of these GenAI tools.** Epic Insights summaries run 16M+ times/month. Per Epic's HIMSS 2026 "Real Results, Right Now" post: more than 200 organizations are actively using Penny for revenue-cycle coding, where "coding-related denials have dropped by more than 20%" and "organizations are creating denial appeal letters 23% faster with Penny." Native MCP-style agentic capabilities were previewed at UGM 2025 ("Sci-Fi" theme). (Fierce Healthcare + 5)
- **Cosmos AI / "Curiosity"**: Epic's foundation models trained on 1,760+ hospitals / 300M patients of de-identified data. (CNBC) (Advisory)

5. Model Context Protocol (MCP) in healthcare — state of the art

- MCP is an open standard launched by **Anthropic in November 2024**; widely described as "USB-C for AI." It solves the M×N problem of connecting N AI systems to M data sources. (Vouched) (Keragon)

- **Production examples in healthcare as of 2026:**
 - **CharmHealth** launched an open-source MCP server in Feb 2026 (github.com/CharmHealth/charm-mcp-server) for EHR data access. [Themomentum](#) [HIT Consultant](#)
 - **FHIR MCP Server** (Momentum) — open-source MCP server exposing FHIR resources via natural language.
 - **EHR-MCP** (Keio University, arXiv 2509.15957, Sept 2025) — real-world deployment showing GPT-4.1 via LangGraph ReAct agent retrieving clinical data through MCP tools at near-perfect accuracy on simple infection-control tasks. [arXiv](#)
 - **Keragon** offers a purpose-built Healthcare MCP exposing 300+ healthcare integrations (Athenahealth, Elation, Healthie) with HIPAA controls at the server layer.
 - **Hearst Health / FDB** publicly endorses MCP as "the missing connective tissue that makes AI in healthcare safe, useful and scalable" (per FDB EVP Charles Tuchinda). [Healthcare IT News](#)
- **Critical compliance note: HIPAA compliance is NOT built into the MCP protocol.** It must be implemented at the server layer: authentication, RBAC, audit logging, encryption, PHI minimization, and BAAs with every downstream model provider. Multi-user authorization (not just OAuth) is the central unresolved challenge — Arcade.dev and others are building this layer.
- **Production gating:** A joint survey of 400+ payer, provider, and pharmaceutical company executives by Bessemer Venture Partners, Bain & Company, and Amazon Web Services (reported by Healthcare Dive, June 2024) found that "just 30% of AI pilots have reached the production stage, as executives struggle with security and data challenges, as well as integration costs and limited in-house expertise." The Philips Future Health Index 2024 (US-specific June 18, 2024 press release) found "88% of US Healthcare Leaders say the use of automation for repetitive tasks is critical for addressing staff shortages." [Healthcare Dive + 2](#)

6. Workflow context — where social workers and care coordinators actually live

- **Make-A-Wish-approved referrers inside a children's hospital are exactly the roles already piloting AI chat tools:** social workers, RN care coordinators, child-life specialists, oncology nurse navigators, and attending physicians. This is the bullseye persona.
- **AI in pediatric care-coordination workflows is nascent but accelerating.** Example: **DAPHNE** (Nationwide Children's Hospital, in pilot RCT) uses Claude/GPT/Mistral/Gemini behind an AWS HIPAA-compliant chatbot to handle health-related social-needs screening across pediatric primary-care clinics — recruiting 10 providers across 14 sites including physicians, nurses, social workers, and care coordinators. [PubMed Central](#) [PubMed Central](#)
- **No current EHR-integrated tool screens patients for charitable-program eligibility (Make-A-Wish, similar orgs).** This is a true white-space.

- **The known pain point — under-referral — is a recognition problem, not a data problem.** All the qualifying signals (age, diagnosis, specialty, treatment history) already exist in the EHR. The work is humans remembering to check.

7. Market and regulatory context

- **HIPAA + BAAs:** Consumer-tier ChatGPT, Claude Pro/Max/Team, Gemini consumer, and Microsoft Copilot Business are **NOT BAA-eligible**. Only sales-managed ChatGPT Enterprise/Edu/Healthcare, Claude Enterprise (HIPAA-ready), and Azure OpenAI / AWS Bedrock / Google Vertex AI with BAA can lawfully process PHI. A signed BAA is necessary but not sufficient — runtime PHI controls, RBAC, audit logging, and de-identification at the prompt boundary are the customer's responsibility. (WitnessAI)
- **ONC HTI-1 final rule (effective Feb 8, 2024):** First federal regulation imposing transparency on AI in certified health IT. Requires developers to disclose 31 source attributes for Predictive Decision Support Interventions (DSIs) and meet the FAVES standard (Fair, Appropriate, Valid, Effective, Safe). ONC-certified health IT is used by **>96% of U.S. hospitals**, so almost any production AI tool touching EHR data is in scope. The rule does not directly regulate generative AI, but the principles (model cards, source attribution, bias testing) will be referenced by hospital procurement. (Mintz + 3)
- **State laws and FTC scrutiny** are increasing on AI handling health data. NIST AI Risk Management Framework and ISO/IEC 42001 (the standard Microsoft Copilot Health was certified to) are becoming buyer checkboxes.
- **Appetite for connecting AI to internal patient data via MCP-style integrations is high but cautious:** the JAMA study showed hospitals with the most robust local evaluation practices for predictive AI moved more slowly to generative AI — i.e., the sophisticated buyers want governance answers first. (Beckers Hospital Review)

8. Similar U.S. wish-granting / pediatric charitable organizations

Organization	Age	Illness Required	Medical Sign-Off	Income Cap	Not
Make-A-Wish America	2½-17	Critical / life-threatening	Yes — treating physician + NMAC	None	16,2 FY. elig
Starlight Children's Foundation	Pediatric (hospitalized)	Any serious illness	No individual sign-off — delivered via hospital partners (824 facilities served in 2024, reaching 3M+ children)	None	Pro by (spe (Hc VR Wa
(Starlight Children's Founda...					

Organization	Age	Illness Required	Medical Sign-Off	Income Cap	Notes
					Th Not
Children's Wish Foundation International	<18	Limited life expectancy	Yes	None	Atl: NC Ch: as c for thru
Kids Wish Network	<18	Life-threatening	Implied (open referral form)	None	Flo Na cha Tar Tin Wi to r Nav Wi con
Sunshine Foundation	3-18	Chronic/lifelong (not life-threatening) — fills the MAW gap	Yes — medical authorization	\$75,000 household Sunshine Foundation Oklahoma ABLE Tech	PA- 197 Su (pro 41,9 Su ~7- Su
Dream Foundation	18+	Life-limiting, ≤12-mo prognosis (≤9 for hospice)	Yes — physician signature mandatory	"Lack resources"	Sar Dr 34,1 sin VA equ

Bottom line: All physician-sign-off models (Make-A-Wish, Sunshine, Dream Foundation, CWFI) share the same screening bottleneck Make-A-Wish has. A clinician-facing MCP screening agent that supports the Make-A-Wish criteria is **easily extensible to Sunshine and Dream Foundation**

with criteria-only changes — multiplying the addressable use case while keeping the integration story identical.

Details

The pitch in plain language

A children's-hospital social worker, oncology nurse coordinator, or child-life specialist opens her Microsoft 365 Copilot, ChatGPT for Healthcare, or Claude Enterprise sidebar — already provisioned by IT, already BAA'd, already used daily for chart summaries and inbox drafting. She types: **"Show me my unit's patients who are likely eligible for a Make-A-Wish referral."** The MCP agent calls into the hospital's FHIR API (Epic, Oracle Health, athenahealth) and the EHR, applies the 12-specialty + age + prior-wish criteria, surfaces a ranked list with rationale, and offers to draft a one-click referral inquiry pre-populated for wish.org. The clinician confirms, the parent is contacted by the local chapter — the only thing the AI replaced was the social worker's mental checklist and a 15-minute form.

Why the timing window is now (~12-18 months)

1. **Hospitals are deploying enterprise GenAI workspaces in 2025-2026 anyway.** OpenAI for Healthcare, Claude for Healthcare, Microsoft Copilot, and Epic Art/Emmie/Penny are buying motions already in flight. An MCP plug-in is additive, not disruptive.
2. **MCP is moving from spec to production.** CharmHealth, Keragon, Hearst Health/FDB, Anthropic's healthcare connectors (Jan 2026), and academic deployments at Keio prove the pattern works inside real hospitals.
3. **Make-A-Wish's own published "1 in 2" reach gap** is the most quoted statistic in the org's fundraising — closing it is mission-aligned, not vendor-pushed.
4. **ONC HTI-1 transparency rules** are setting the bar; an MCP-based screening agent that is fully auditable (every patient flag explainable, every API call logged) is easier to certify than opaque predictive models.

Risks and what would change the recommendation

- **A hospital's Privacy & Security Office may refuse to expose patient cohorts to even a BAA-covered LLM** until ONC HTI-1 v2 (HTI-2) and FAVES enforcement clarify. Mitigation: ship a de-identified-screening + clinician-confirmation flow rather than auto-referral.
- **Make-A-Wish itself may prefer to control the referral funnel** to avoid duplicate or unsuitable referrals; partnership with Make-A-Wish America's National Office (Phoenix HQ) and the National Medical Advisory Council is essential to legitimize the tool.

- **Chapter fragmentation:** 58 U.S. chapters operate semi-autonomously with their own medical advisors. A pilot must be sanctioned both nationally and at the chapter the participating hospital sits in. (Wikipedia) (Wish)
 - **Wish-once-per-child rule** is hard data — the tool must check Make-A-Wish's referral database and ideally the few credible peer-org databases (Sunshine, Dream Foundation, CWFI) to avoid duplicate referrals. This is a real data-sharing problem and a possible defensible moat. (Make-A-Wish)
-

Recommendations

Staged plan, with thresholds that should change the call:

Stage 1 — Validate with one anchor partner (0–3 months)

- **Target a single named pediatric academic medical center already deploying ChatGPT for Healthcare or Claude Enterprise.** First-call list, in priority order: **Boston Children's** (OpenAI launch partner, John Brownstein public sponsor of AI experimentation), **Stanford Medicine Children's Health** (Nigam Shah, ChatEHR precedent), **CHOP** (Abridge adoption), **Texas Children's** (conversational AI precedent).
- Co-design with **one pediatric oncology social-work team** plus the **chapter Medical Advisory Council member** for that geography. Make-A-Wish national co-sign is required — request a meeting with the **National Medical Advisory Council** and the Medical Outreach & Intake team in Phoenix.
- **Kill the project if no anchor hospital will commit to a 10-week paid pilot** by month 3 — the business case depends on hospital procurement appetite, not just clinician enthusiasm.

Stage 2 — Build the minimum-viable MCP tool (3–6 months)

- **Architecture:** MCP server exposes 4 tools — `(list_eligible_patients(unit, date_range))`, `(get_patient_eligibility_rationale(patient_id))`, `(check_prior_wish_status(patient_id))`, `(draft_referral(patient_id, contact_info))`. Server runs in the hospital's cloud (Azure / AWS / GCP) under the hospital's BAA. No PHI leaves the hospital tenant.
- **Models supported on Day 1:** Claude (via HIPAA-ready Enterprise or AWS Bedrock), ChatGPT for Healthcare (via API + BAA), Microsoft Copilot via Azure OpenAI. Skip consumer tiers entirely.
- **Compliance bar:** SOC 2 Type II, HIPAA risk analysis, ISO/IEC 42001 (the standard Microsoft Copilot Health holds), and an HTI-1-style model card published from Day 1.
- **Threshold to expand:** $\geq 30\%$ of pilot-team-identified candidates result in submitted referrals AND ≥ 3 minutes saved per screening event. Below either, redesign.

Stage 3 — Multi-org expansion (6-12 months)

- Extend criteria packs to **Sunshine Foundation** (chronic-condition gap) and **Dream Foundation** (AYA patients aging out of MAW at 18). Same MCP server, new eligibility rule sets, near-zero marginal integration cost.
- **Begin discussions with Make-A-Wish America HQ about a national MOU** — proof points from Stage 2 are required first.

Stage 4 — Defensible moat (12-18 months)

- **Build a federated cross-organization "wish-once" registry** so the tool can prevent duplicate referrals across MAW, Sunshine, and Dream Foundation. This is where regulatory friction is highest and incumbents have not built.
- **Publish the MCP eligibility criteria pack open-source**, but keep the orchestration, audit, and registry proprietary.

Benchmarks that would change the recommendation

- **Pivot away if:** (a) Make-A-Wish HQ explicitly objects to algorithmic pre-screening; (b) HTI-2 / FDA explicitly classifies eligibility-screening agents as SaMD requiring 510(k); (c) Epic announces a native Art-agent skill that does this for free.
 - **Accelerate if:** (a) one Top-10 children's hospital signs a paid pilot; (b) Make-A-Wish national co-promotes the tool; (c) Anthropic, OpenAI, or Microsoft promotes the integration in their healthcare partner marketplaces.
-

Caveats

- **The "1 in 2 eligible children" reach rate** is repeatedly quoted on wish.org/wishes-forever but the underlying methodology (2018 figure of 15,600 wishes vs. 27,000 critically-ill diagnoses) has not been formally updated; the FY24 framing ("for every wish we fulfill, two more are waiting") implies the gap has widened, but Make-A-Wish has not published a current rate. The order of magnitude (a large under-served population) is well-supported; the exact denominator is not. [Wish](#)
- **National diagnosis breakdown for FY24 wishes is not published** in the FY24 Annual Report. The most-cited figure (~45% oncology) is from a 2018 NPR interview with then-MAW VP of Mission Delivery Kim Sandys, plus the Nationwide Children's Hospital cohort study (45% cancer, 16.5% blood/kidney, 14.5% neurological, 14% genetic). Verify before quoting in a customer-facing pitch. [NPR](#)

- **Generative-AI adoption statistics conflate categories.** "Using generative AI" can mean anything from one pilot inbox-drafting use case to enterprise-wide deployment. The JAMA 31.5% figure refers specifically to EHR-integrated GenAI; the McKinsey 50% figure includes any production GenAI use case. Both are directionally consistent — adoption is real and material — but quoting them interchangeably misrepresents the data.
- **MCP standards are evolving fast.** As of mid-2026, multi-user authorization for healthcare MCP is still being solved by vendors like Arcade.dev. Treat the MCP architecture choice as a directional bet, not a finalized standard. Build the agent so the underlying transport can be swapped without re-litigating eligibility logic.
- **Make-A-Wish operates through 58 semi-autonomous U.S. chapters.** National-office buy-in does not guarantee chapter cooperation, and chapter-level Medical Advisory Council members will be the local gatekeepers in any pilot.
- **CharityWatch's negative ratings on Children's Wish Foundation International and Kids Wish Network** are material — do not position these as peer reference customers in a pitch. Use Sunshine Foundation and Dream Foundation as the credible "extensible to" stories instead.
- **Future-tense claims in vendor marketing** (e.g., "agentic AI will transform prior authorization") should be discounted; only claims attached to a named hospital and a published metric belong in a customer pitch.